

P8133 Term Paper: N-of-1 Trial Simulation

Adina Zhang & Christian Pascual

12/16/2019

1 Introduction

1.1 Clinical Relevance

In the age of personalized medicine, N-of-1 trials show great promise in determining relevant and effective treatments in individual patients [1]. Randomized control trials (RCTs) are considered the gold standard for investigating treatment effects, but the results of these studies are limited to average treatment effects which may not be applicable to individual patients. RCT can be prohibitively expensive, so relevant RCT data may not exist or are not easily accessible for patients who may be ineligible or unable to participate. These weaknesses highlight how the clinical trial landscape would benefit from N-of-1 studies.

N-of-1 trials are multiple crossover trials where a patient periodically switches between treatments. They are characterized by the ability to calculate individual treatment effect (ITE) while still maintaining balanced or randomized treatment assignment. The goal of an N-of-1 trial is to determine effective treatments for individuals while minimizing adverse events and maximizing the patient benefit. Patients are encouraged to contribute feedback and tailor the outcomes of interest to their own personal needs, giving N-of-1 trials a unique ability to engage patients in their own medical decisions. For all their strengths, N-of-1 trials are best suited to a particular disease type. Diseases that are chronic or slowly progressing are most suitable because acute clinical and environmental factors may bias results resulting in incorrect conclusions. Clearly symptomatic diseases or defined biomarkers are necessary to establish measurable outcomes and judge the efficacy of treatment.

The design of an N-of-1 trial is dictated by many design and analytical factors, as well as treatment characteristics. All of which impact the ability to select the best treatment for individual patients and correctly identify ITEs. Determining the order of treatment assignments, considering short and long-term time trends, and accounting for treatment carryover and run-in effects are just a few of the considerations a researcher must account for. The sample size of an N-of-1 trial is based purely on the design of the trial, which has implications on both power and patient attrition. To limit the scope of our paper, we focus on design and treatment characteristics.

1.2 Problem Formulation

The single-subject nature of N-of-1 trials lends them to particular design issues which have the potential to affect statistical and clinical endpoints. Our paper is inspired by a simulation study by Percha et al, and we plan to simulate how tuning design parameters affect downstream analyses [2]. We are also interested in investigating how varying treatment parameters affect design considerations, given heterogeneous patient response to treatment. To carry out this investigation, we do a two-fold investigation. First, we present a case study to motivate how different design parameters affect

statistical and clinical endpoints, such as power and ITE estimation. Second, we perform more simulations varying treatment-related parameters and see how these effect the same end points.

2 Methods

2.1 Simulation Software

Our simulation software was designed to allow a user to vary multiple parameters to generate N-of-1 trial data for many different clinical scenarios and treatment types. We took inspiration from the simulation work done by Percha et. al and based many of our simulation decisions off of their paper. The simulation software was created in R and is hosted in a Github repo.

2.2 Simulation Parameters

The parameters of our software can be divided into 3 categories: design parameters, treatment parameters and outcome parameters.

Design parameters are parameters that a researcher or clinician would actually have control of in the creation of an N-of-1 trial, which include: 1) sampling frequency within a period, 2) the period length of a single treatment, 3) the number of treatments to be used in the trial and 4) the number of blocks of treatment. These parameters dictate how long the N-of-1 trial will be, which ultimately decides the final sample size of the trial itself.

Treatment parameters are parameters that concern characteristics of the treatments to be used during the course of the trial. We modeled four specific aspects of a treatment that a user can control. The first treatment parameter is the treatment effect, which captures how much hypothetical change a treatment will exert on a clinical outcome of interest. The next two treatment parameters simulate the effects of a patient stopping and starting treatment. Each treatment has a carryover effect, which is how long it takes for it to completely lose its effect after cessation. Conversely, each treatment also has a run-in effect, which details how long it takes for a treatment to reach its full effect. Finally, each treatment is associated with a noise parameter that represents deviations that a treatment may have from its true treatment effect. Carryover roughly models how the body removes compounds from the body, while run-in models the time it takes for treatments to bind to their specific targets in the body. Treatment noise models how environmental factors such as stress may affect the body's potential to respond to treatment.

Outcome parameters relate to the specific clinical endpoint that a subject is interested in intervening upon, such as blood pressure. For the sake of limiting scope, we focused on a continuous outcome. We refer to the outcome unaffected by any treatment effect as the *baseline value*. We modeled a subject's baseline value through time as a Markov process. Users can specify a starting value and a baseline noise value that dictates how far the baseline value can deviate from the last baseline observation. These parameters attempt to capture the natural variability in many clinical endpoints that may be encountered in N-of-1 trials.

Table 1: Tunable parameters in our N-of-1 trial simulation software

Parameter	Notation	Type	Description
Sampling Frequency	F	Design	Number of times patient is sampled in one period
Period Length	P	Design	Length of the treatment period
Number of Treatments	T	Design	How many treatments used in trial
Treatment Order	O	Design	Order of treatments given to patient in single block
Number of Blocks	B	Design	How many treatment blocks used in trial
Treatment Effect	E_k	Treatment	Effect of treatment T on the baseline
Carryover	τ_k	Treatment	Constant affecting how long it takes for effect of treatment k to go away
Run-in	γ_k	Treatment	Constant affecting how long it takes for full effect of treatment k to occur
Treatment Noise	σ_k^2	Treatment	How much will the treatment k effect vary overall
Baseline Start	μ_b	Outcome	Starting value for the baseline Markov chain
Baseline Noise	σ_b^2	Outcome	How much can the Markov chain move at each time point

2.3 Data Generation Process

Table 1 summarizes all of the parameters that a user can tune in our simulation software. We denote $B(t)$ as the baseline outcome value of a subject at time t , which has values going from $1, \dots, F \times T \times P \times B$. A subject's baseline outcome value is modeled as a Markov process. Given a starting value μ_b , the next value in the chain is calculated using σ_b^2 as:

$$\mu_{t+1} = \mu_t + M, \quad M \sim N(0, \sigma_b^2)$$

We denote the treatment effect for treatment k as $X(t)$ which has two subcomponents: a deterministic component $X_{det}(t)$ and a random component $X_{ran}(t)$. The treatment is modeled as an exponential decay from a starting value a to a target value b . Using α to denote a general time constant, the decay was modeled as:

$$X_{det}(t) = b + (a - b)e^{-\Delta t/\alpha}$$

where Δt is a vector of values spanning an entire period length. When treatment is currently being taken, $\alpha = \gamma_k$ since its full effect is running in, and the start and target values are 0 and E_k respectively. When the treatment is not being taken, $\alpha = \tau_k$, and the start and target values are E_k and 0. The random component is a normal random variable with variance σ_k^2 , so $X_{ran}(t) \sim N(0, \sigma_k^2)$. Each treatment has both a deterministic and random component, so the cumulative effect over all T treatments is represented as the total sum over all T treatments.

$$\sum_{i=1}^T X^i(t) = \sum_{i=1}^T \left(X_{det}^i(t) + X_{ran}^i(t) \right)$$

We denote $Y(t)$ as the *observed* clinical value for a subject, and it is simply the sum of the baseline outcome value and the total treatment effect.

$$Y(t) = B(t) + \sum_{i=1}^T X^i(t)$$



Figure 1: An example simulation output. Baseline value was initialized at 100 with $\sigma_b^2 = 5$. Treatment A (red) has a true effect of -30 with a carryover constant of 2, run-in constant of 2 and $\sigma_A^2 = 1$. Treatment B (blue) is a placebo with no treatment effect 0, no carryover or run-in and $\sigma_B^2 = 1$. Other treatments (C) were not specified and do not contribute to the observed effect.

2.4 Simulation Scope

For the purposes of limiting the scope of term paper, we have chosen not to model particular important aspects of an N-of-1 trial. We have chosen not to model potential interactions between treatments and long-term time trends in the clinical outcome that are commonly seen in trials. Some trials incorporate wash-out periods to account for run-in, but this was not included in our software. Since we are focusing on design and treatment parameters, we have not modeled patient-specific autocorrelation in the simulation process.

2.5 Statistical Analyses

For our analyses, we have chosen to perform regression analyses as recommended by Kravitz et. al in the AHRQ Handbook [1]. The clinical outcome will be modeled as a linear function of $T - 1$ indicator variables for each treatment using a placebo as the reference and $B - 1$ indicator variables for each block using the first block as the reference. The model can be written as:

$$Y_{kb} = \beta_0 + \beta_1' \mathbf{T} + \beta_2' \mathbf{B} + \epsilon_{kb}$$

where \mathbf{T} is the matrix of treatment indicator variables indexed by k and \mathbf{B} is the matrix of block indicators indexed by b . After some exploratory analytical testing, we found this model to provide accurate estimates while remaining estimable. We attempted to use linear mixed-effects models, but we found that we were not generating enough data for the models to converge.

A significant result in a simulation is defined as when the regression coefficient associated with the active treatment to be significant ($p < 0.05$). Overall power of a given set of simulations was defined as the proportion of which obtain a significant result for the active treatment. Finally, we define the probability of correct selection (PCS) as the proportion of simulation trials where the resulting coefficient estimate for treatment satisfies a pre-specified condition, such as maximization of effect or choosing the treatment that avoids unnecessary under- or over-treatment.

For our study, we investigate how different sets of design and treatment parameters have the potential to affect the overall power, effect estimation and PCS. We use this criteria as a balance between researchers who may want to maximize power and accurate estimation against patient interests, who will want to find the correct treatment as soon as possible while minimizing harm and cost.

2.6 Performance Metrics

Our case study highlights just one use case for N-of-1 trials. For our paper, we chose to use three performance metrics to explore how different parameters affect the needs of a trial: average power, average deviation and probability of correct selection (PCS). All performance metrics were calculated from the results of 100 simulations, given a set of treatment parameters. For power, a linear model was created after each simulation, and the significance of the treatment was checked. Average power was calculated as the proportion of the 100 simulations that found the treatment to be significant. For average deviation, the deviation away from the (known) true effect was calculated from the model estimate. Average deviation was calculated as the mean deviation of the 100 simulations. PCS was calculated under a particular selection criteria. For our paper, we explored two different criteria for correct selection: 1) correct when choosing the treatment with maximal effect and 2) correct when choosing the treatment that most accurately estimates an effect within a given window of acceptable effect. PCS was calculated as the proportion of simulations that correctly select the known optimal treatment.

3 Results

3.1 Case Study: Active Treatment vs Placebo

To demonstrate the nuances of design consideration, we present a case study using our simulations. We assume two treatments with a continuous outcome measurement: *A* has a treatment effect of -3 with 2 run-in and 2 carryover days. *B* is the placebo with no treatment effect and has 1 run-in and 1 carryover day (Figure 2). Our choice of tuning design parameters resulted in the following simulation scenarios:

1. Treatment selection: *ABAB*, *ABBA*, *BABA*, *BAAB* holding sampling frequency to 1, period length to 5 days, and the number of blocks to 2
2. Sampling frequency: 1 to 30 times in a period holding treatment order to *ABAB*, period length to 5 days, and the number of blocks to 2
3. Period length: 2-120 days holding treatment order to *ABAB*, sampling frequency to 1 per period, and the number of blocks to 2
4. Number of blocks: 1-6 blocks holding treatment order to *ABAB*, sampling frequency to 1 per period, and period length to 5 days

Figure 2 was derived from simulating each scenario and parameter 50 times. Effect estimates were calculated from regression adjusting for block number as described previously. The median value of effect estimates were calculated from statistically significant regression estimates and used for comparison.

From Figure 2, the treatment order *ABBA* has the closest estimate of treatment effect with a median of -2.45 and variance of 0.18. The differences in treatment effect are likely a result of the carryover and run-in effects. Given that our simulation assumed 5 day periods and sampled once a period, carryover and run-in effects become prominent.

Increasing sampling frequency results in higher accuracy of treatment effect with decreasing variance. However, from Figure 2, the increased sampling frequency has diminishing results and the best estimates started to plateau. Around 5 samples in a 5 day period results in a median treatment effect of -2.83 with a 0.04 variance while 10 samples in a 5 day period results in a median treatment effect of -2.91 with a 0.02 variance. Sampling can be invasive to a patient but critical to power of the study. This simulation demonstrates that it is possible to identify a lower sampling frequency without compromising the best estimate of treatment effect.

Like sampling frequency, increasing period length increases accuracy of treatment effect estimation over time along with decreasing variance. From Figure 2, treatment effect plateaus between period lengths of 15 to 30 days, where median treatment effects are -2.70 and -2.90 respectively. Again, large period lengths can be exhaustive to a patient, especially if there are negative side effects involved. Minimizing period length while still maintaining power and effect estimate are important steps to consider. It should also be noted from Figure 2 that the results when period length is 2 days are not significant. This is probably a result of low power in addition to carryover and run-in effects. Effects from treatment A and B likely run into each other within the span of two days.

We varied block length and in Figure 2 we observe no significant difference in effect estimate with increasing blocks. However, variance does decrease with increased blocks. We surmise that the reason we do not observe any differences is that our simulation does not model time-trends. We do observe variance decreased because sample size increases with blocks thereby increasing certainty in the estimates.

3.2 Sample Size & Power

Figure 3a is the result of a series of simulations designed to investigate the effect of changes to treatment effect sizes on average power and effect estimation. We kept the number of blocks constant at 2 and the treatment order at *ABBA* (*A* being the active treatment, *B* being placebo). The initial baseline value was 0, and baseline noise was set to 1. For these simulations, the effect of the active treatment was a linear function of the baseline noise. The active treatment was given carryover and run-in constants of 2 and a negligible noise. Keeping sampling frequency at 1/period and looking at different period lengths from 1 to 20 days, we looked at the average power for a set of 100 simulations. We have a supplemental table at the end laying out the simulation parameters used for each figure.

The result is a trend line that shows the relationship between the final sample size the resulting power from that sample size, stratified by different effect sizes. We observed that for extremely small effect sizes (< 0.5) that trials must be extremely long to detect a meaningful effect. For the largest effect size we simulated, we observed that a sample size of about 80 measurements was needed to

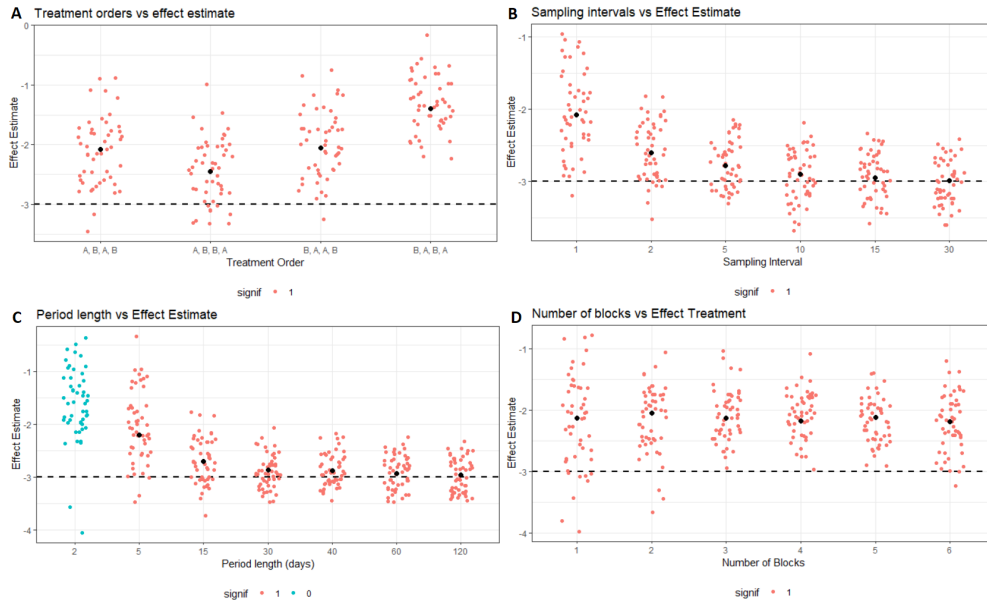


Figure 2: Simulation results on effect estimate after tuning for design parameters, including (A) treatment orders, (B) sampling frequency, (C) period length, (D) number of blocks. Each tuning parameter was run with 50 simulations. Drug A has a true effect of -3 while Drug B is a placebo with 0 effect. Black points are median effect sizes among statistically significant results.

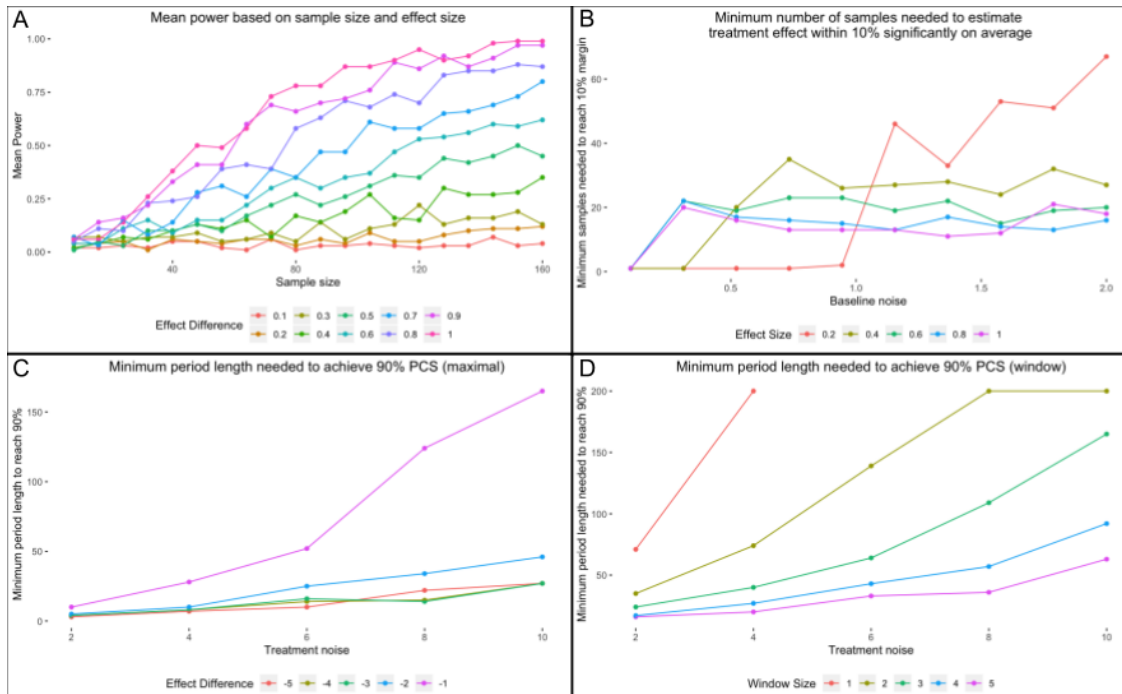


Figure 3: Simulations showing how varying treatment parameters affects power, treatment effect estimation and PCS.

reach 80% and about 110 to reach 90% power. This equates to a period length of about 10.

3.3 Minimum Sample Size & Average Deviation

Figure 3b shows the result of simulations designed to look at the minimum sample size needed to accurately estimate the true treatment effect on average. “Accurate estimation” was defined as the mean deviation between the estimated treatment effect and the true effect across 100 simulations. The design and treatment parameters were kept the same as in the sample size and power simulations. Keeping sampling frequency at 1, period length was increased until the average deviation was within 10% of the true effect. This process was repeated for different effect sizes. We found that a sample size of about 50 was enough to accurately estimate the treatment effect, even as baseline noise increased. However, for extremely small effects, the necessary sample size for accurate estimation jumped to about 100.

3.4 Selection for maximal effect

Figure 3c looks at simulations done to investigate how well an N-of-1 trial is able to correctly choose between two competing treatments. Out of 100 simulations for a given a combination of treatment noise and difference in treatment effects, we calculated the PCS. The design parameters were kept the same as in 2a and 2b, except for treatment order. A placebo was used as a third treatment, and the order was the two active treatments followed by the placebo. The optimal active treatment had a reducing effect of 5. Both active treatments had run-in and carryover constants of 2. We observe that for small absolute differences in the competing treatment effects that especially large sample sizes are needed to detect the effect. For example, an effect size of 1 corresponds to a period length of 33 is needed to correctly choose the best treatment 90% of the time. As the effect difference increases, the necessary sample size to achieve a PCS of 90% seems to level off at about a period length of 20.

3.5 Selection in an optimal treatment window

Figure 3d examines a different definition of probability of correct selection. For these simulations, there were 3 competing active treatments, accompanied by a placebo. One treatment was considered optimal, while the other two had treatment effects that represented both under- and over-treatment. We investigated the minimum sample size required to correctly identify the optimal treatment 90% of the time over 100 simulations. The optimal treatment effect was known, and a treatment was selected based on the smallest deviation from this known effect. Sampling frequency was kept at 1 per period, and the treatment order was set at: optimal, under-treatment, over-treatment, then placebo. All active treatments had the same noise, carryover and run-in. The initial baseline parameters were the same for the previous treatment selection simulation.

We observed that for small windows of efficacy (effect differences $\pm 1, 2$) the minimum sample size necessary to pick out the optimal treatment grows roughly exponentially as the treatment noise increases. We chose to limit the number of candidate sample sizes to 200, which we see is reached by the small window sizes. For more generous windows, the exponential increase is slowed. For an effect size of 1 ($\sigma_t^2 = 5$), we see that a minimum sample size of about 50 observations is needed

to correctly select the best treatment 90% of the time. From the treatment order and number of blocks, this translates to a period length of about 7 days per treatment.

4 Discussion

4.1 Study Intent

The purpose of our study was to use simulation to investigate how different treatment parameters have the potential to affect important statistical and clinical endpoints, which have the potential to be at odds with each other. A statistician may desire the extension of a study for the purposes of power and accurate estimation, but this may come at the cost of higher patient attrition or extended exposure to treatment side-effects. Simulation is a powerful tool for allowing researchers to better plan out their designs in hopes of striking a balance between these two realities, but as we have discovered over the course of this project, this benefit is only as good as the realism of the simulation itself. We hoped to extend the work of Percha et. al to treatment selection and non-ideal cases where carryover and run-in effects are present.

4.2 Recommendations

4.2.1 Case Study

Within our simulations, we confirm that treatment order, sampling frequency, and period length to be important design considerations. Had our simulation included time-trends, we expect block length to be critical as well. We noticed the simulations were strongly influenced by wash-in and carryover effects, so we would recommend considering washout periods between treatments or consider sampling after treatment effects are stabilized to decrease these effects. Period length and sampling frequency should be balanced in a way that allows for accurate treatment selection without making the trial unnecessarily long. Our case study demonstrates the importance of simulations in their ability to decompose complex design and environmental factors to better understand how analysis can be maximized without compromising the patient’s healthcare.

4.2.2 Power

The length of the N-of-1 trial is the sole determinant of the sample size, and researchers are able to change their sample size by tuning the design parameters. From our simulation findings, we observe that N-of-1 trials run a high risk of being highly underpowered, especially in cases where the active treatment does not do much better than the placebo or when the clinical outcome of interest is highly variable. It is no surprise that researchers must try to make as many observations as possible to achieve acceptable power, but it is helpful to keep a sort of minimum necessary sample size to have some confidence in truly detecting effective treatments. For larger effects, we found that sample sizes of about 80 observations were enough to start approaching 80% power. With this in mind, researchers designing N-of-1 trials should be prepared to conduct studies spanning 3-4 months, assuming one observation per day. Shorter trials may still be effective at accurately estimating the treatment effect, but these effects may not be truly significant.

4.2.3 Probability of Correct Selection

Our simulations find that the definition of “correct” treatment selection has critical implications on the necessary sample size needed to identify the correct treatment for a patient. In simple cases where we are only seeking the treatment with maximal effect, period length of about 50 seem to suffice in correctly identifying the maximal treatment effect, even in cases where the treatment noise is large.

This is not the case when we must avoid both under-treatment and over-treatment. This problem may arise in cases where an adequate treatment is known, but the correct dose to use is not. Compared to the corresponding line in Figure 3c, having to account for a window necessitates a greater period length to avoid under- or over-treatment. This difference in sample size is even greater when the exact window for optimal treatment is small. For example, even for effect sizes of 2 ($\sigma_t^2 = 2.5$), the necessary period length to identify a maximal treatment effect is about 25, compared to about 150 needed to identify it in a window. This finding highlights the fact that including more treatments necessitates greater sample size to accurately estimate each of their effects. Researchers should carefully define their goal in correctly selecting a treatment for a subject since it can greatly alter the resulting design that is needed to achieve it.

4.2.4 Effect Estimation

In terms of the three performance metrics we chose to examine, accurate treatment effect estimation may not hold as much importance in an N-of-1 trial. Despite this, we still feel that accurate estimation of treatment effects has critical implications in proper treatment selection. As mentioned before, our simulations found that N-of-1 trials tend to estimate treatment effects accurately faster than they do achieve statistical power. This finding might be useful in cases where it is unknown how a subject will react to treatment. To start estimating the true treatment effect, sample sizes of about 20 observations are needed. This information might be useful to researchers to motivate a reevaluation of the trial design after the effects are estimated, much like an internal pilot.

These findings highlight a looming logistical problem with N-of-1 trials. More reasonable trial lengths from a patient perspective are more likely than not to be underpowered and have biased estimates of the treatment effect. Researchers should combine simulation use and patient consultation to work towards a compromise that fits both their and patient needs. Our simulations also demonstrate that acceptable performance metrics are met when the treatment effects are similar in magnitude to baseline or treatment noise; in cases where these effects are small, researchers may want to discuss the pros and cons of extended treatment regimens for effects that may not be clinically relevant for the patient.

4.3 Limitations

For our project, we made several simulation design choices to limit the scope of our project. We chose to emulate much of our simulation after the Percha study, but we made some changes based on time constraints. We model a baseline clinical outcome as a Markov process to represent a subject outcome of interest in the absence of any treatment. We ignore any longer time trends such as drift or cyclicity, which are apparent in many important biological phenomena. Furthermore, in our simulations, we have only chosen a small subset of parameters in a vast parameter space, so our

findings may differ greatly in other contexts. For example, long carryover and run-in have immediate consequences on proper treatment selection, which may motivate a researcher to incorporate a washout period.

4.4 Group Reflection

This project gave us an increased appreciation for the role of simulation in guiding research. We've been exposed to the use of simulation in evaluating statistical techniques in Prof. Wei's Topics in Computing Class, but this is the first time we've seen it used in a clinical trial context. Although we were unable to incorporate more Bayesian or adaptive techniques into our analytic plan, we feel that the framework we laid out can easily be extended to accommodate this type of analysis. Although we're proud of the work done here, there's a lot of room for refinement. In the end, we both have appreciated the opportunity to explore the N-of-1 studies as an alternative to other study-designs we've been exposed to. They are dizzyingly complex and each individual study will have different aspects to take into consideration, but at least this is where simulation can help to tease through these difficult decisions.

5 References

1. Kravitz RL, Duan N, eds, and the DEcIDE Methods Center N-of-1 Guidance Panel (Duan N, Eslick I, Gabler NB, Kaplan HC, Kravitz RL, Larson EB, Pace WD, Schmid CH, Sim I, Vohra S). Design and Implementation of N-of-1 Trials: A User's Guide. AHRQ Publication No. 13(14)-EHC122-EF. Rockville, MD: Agency for Healthcare Research and Quality; January 2014.
2. Percha B, Baskerville EB, Johnson M, Dudley JT, Zimmerman N. Designing Robust N-of-1 Studies for Precision Medicine: Simulation Study and Design Recommendations. J Med Internet Res 2019;21(4):e12641

Table 2: (Supplement)Parameters used in different simulations for Figure 3

Parameter	Notation	2a	2b	2c	2d
Sampling Frequency	F	1	1	1	1
Period Length	P	Varies	Varies	Varies	Varies
Number of Treatments	T	2	2	3	4
Treatment Order	O	A, B, B, A	A, B, B, A	A, B, C	A, B, C, D
Number of Blocks	B	2	2	2	2
Treatment Effect	E_k	Varies	Varies	-5, Varies	-5, Others Varies
Carryover	τ_k	2	2	2	2
Run-in	γ_k	2	2	2	2
Treatment Noise	σ_k^2	0.01	0.01	Varies	Varies
Baseline Start	μ_b	0	0	100	100
Baseline Noise	σ_b^2	1	Varies	1	1
Number simulations	N	100	100	100	100